

Accessing Covid-on-the-Web dataset in R

This notebook demonstrates how to access and query Covid-on-the-Web RDF dataset. The dataset describes named entities identified in the scholarly articles of the COVID-19 Open Research Dataset (CORD-19), a resource of over 47,000 articles about COVID-19 and the coronavirus family of viruses.

The full description of the dataset can be found here: <https://github.com/Wimmics/CovidOnTheWeb>

You can query the dataset from our Virtuoso endpoint: <https://covidontheweb.inria.fr/sparql>.

Cite this work

When including Covid-on-the-Web data in a publication or redistribution, please cite the dataset as follows:

Wimmics Research Team. Covid-on-the-Web Dataset. University Côte d'Azur, Inria, CNRS. 2020. Retrieved from <https://github.com/Wimmics/CovidOnTheWeb>.

```
library(SPARQL)
```

```
endpoint <- "https://covidontheweb.inria.fr/sparql"
options <- NULL
```

```
prefix <- c('covid','<http://ns.inria.fr/covid19/>',
            'wd',   '<http://www.wikidata.org/entity/>',
            'wdt',  '<http://www.wikidata.org/prop/direct/>')
```

```
sparql_prefix <- "
PREFIX rdfs:  <http://www.w3.org/2000/01/rdf-schema#>
PREFIX owl: <http://www.w3.org/2002/07/owl#>
PREFIX xsd:   <http://www.w3.org/2001/XMLSchema#>
```

```
PREFIX bibo:  <http://purl.org/ontology/bibo/>
PREFIX dce:   <http://purl.org/dc/elements/1.1/>
PREFIX dct:   <http://purl.org/dc/terms/>
PREFIX fabio: <http://purl.org/spar/fabio/>
PREFIX foaf:  <http://xmlns.com/foaf/0.1/>
PREFIX frbr:  <http://purl.org/vocab/frbr/core#>
PREFIX oa:    <http://www.w3.org/ns/oa#>
PREFIX prov:  <http://www.w3.org/ns/prov#>
PREFIX schema: <http://schema.org/>
```

```
prefix wd:    <http://www.wikidata.org/entity/>
prefix wdt:    <http://www.wikidata.org/prop/direct/>
```

```
PREFIX covid:  <http://ns.inria.fr/covid19/>
PREFIX covidpr: <http://ns.inria.fr/covid19/property/>
"
```

Working with article metadata

Query dataset for the articles that mention the term *coronavirus* in their abstracts.

```

query <- '
SELECT (group_concat(distinct ?name,"; ") AS ?authors)
       ?title
       (year(?date) as ?year)
       ?pub
       ?url

WHERE {
  graph <http://ns.inria.fr/covid19/graph/articles>
  {
    ?doc a ?t;
        dce:creator ?name;
        dct:title ?title;
        schema:publication ?pub;
        schema:url ?url;
        dct:abstract [ rdf:value ?abs ].

    optional { ?doc dct:issued ?date }
    filter contains(?abs, "coronavirus")
  }
}
group by ?doc ?title ?date ?pub ?url
order by desc(?date)

'

query <- paste(sparql_prefix, query)

```

```

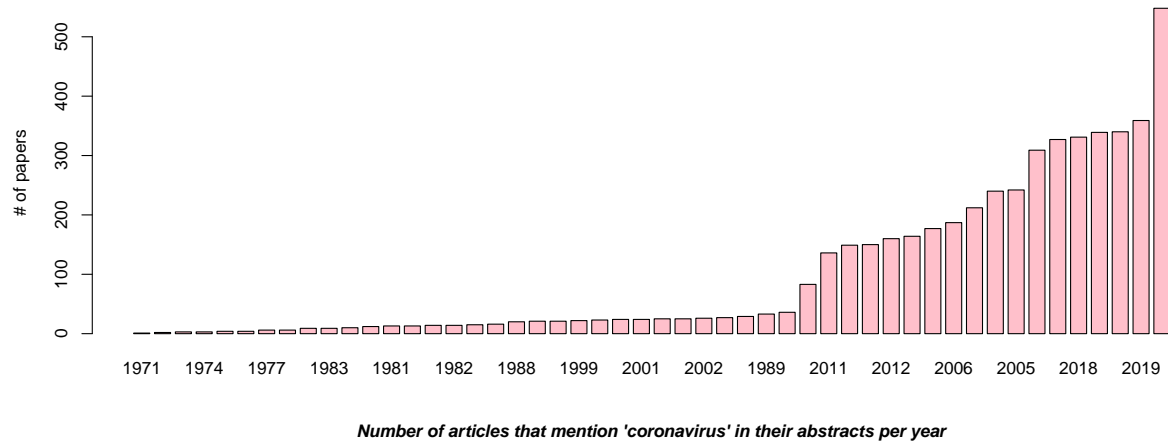
res <- SPARQL(url= endpoint,
              query = query,
              ns=prefix,
              extra=NULL)$results

```

```

barplot(sort(table(res$year)),
        col="pink",
        space=0.3,
        ylab="# of papers",
        sub="Number of articles that mention 'coronavirus' in their abstracts per year",
        font.sub=4)

```



Working with article annotations

Query dataset for the articles referencing *coronavirus* and forms of *cancer* at the same time

```
query_corona_vs_cancer = '
# wdt:P279 = subclass of
# wdt:P31 = instance of
# wd:Q12078 = cancer
# wd:Q1134583 = coronavirus family

select distinct ?article ?dis1 ?dis1Label ?dis2 ?dis2Label #?dis2Subject

from <http://ns.inria.fr/covid19/graph/entityfishing>
from named <http://ns.inria.fr/covid19/graph/wikidata-named-entities>

where {
    # Look for 2 annotations of the same article with Wikidata URIs ?dis1 and ?dis2

    ?annot1 schema:about ?article; oa:hasBody ?dis1.
    ?annot2 schema:about ?article; oa:hasBody ?dis2.

    graph <http://ns.inria.fr/covid19/graph/wikidata-named-entities>
    {
        ?entity1 rdfs:label "cancer"@en. # ?entity1 is wd:Q12078

        { ?dis1 rdfs:label ?dis1Label.
          filter (?dis1 = ?entity1) } # ?dis1 is "cancer"

        UNION

        { ?dis1 wdt:P279 ?entity1;
          rdfs:label ?dis1Label. } # ?dis1 is a subclass of "cancer" (at any depth)

        UNION

        { ?dis1 wdt:P31 ?entity1;
```

```

        rdfs:label ?dis1Label. } # ?dis1 is an instance of "cancer" or a subclass

?entity2 rdfs:label "Coronaviridae"@en. # ?entity2 is wd:Q1134583

{ ?dis2 rdfs:label ?dis2Label.
filter (?dis2 = ?entity2) }

UNION

{ ?dis2 wdt:P279 ?entity2;
  rdfs:label ?dis2Label. } # ?dis2 is a subclass of "Coronaviridae" (at any depth)

UNION

{ ?dis2 wdt:P31 ?entity2;
  rdfs:label ?dis2Label. } # ?dis2 is an instance of "Coronaviridae" or a subclass

}

}
order by ?dis1 ?dis2
limit 1000
'

```

```

query_corona_vs_cancer <- paste(sparql_prefix, query_corona_vs_cancer)

```

```

res <- SPARQL(url= endpoint,
              query = query_corona_vs_cancer,
              ns=prefix,
              extra=NULL)$results

```

```

# remove label decorations
res <- data.frame(lapply(res, function(x) {
  gsub("@en", "", x)
}))
res <- data.frame(lapply(res, function(x) {
  gsub("\\\"", "", x)
}))

```

Visualize query results in different ways

Plot hierarchically-clustered heatmap

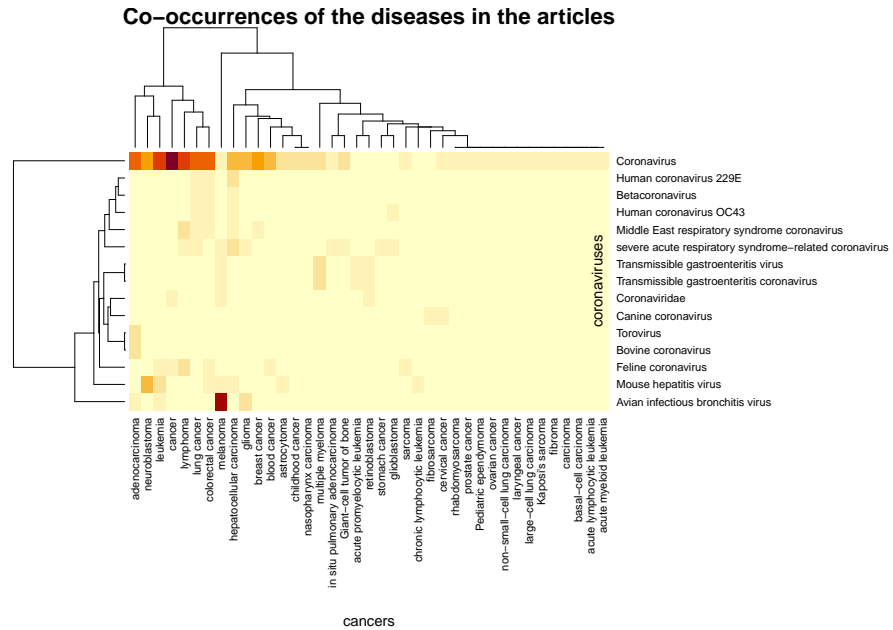
```

hm <- table(res[, c('dis2Label','dis1Label')])
hm <- as.matrix(hm)

heatmap(hm, margins = c(15,0), #Rowv = NA, Colv = NA,
        col = hcl.colors(12, "YlOrRd", rev = TRUE),
        scale="none",
        cexCol=0.9,
        cexRow = 0.9,
        main="Co-occurrences of the diseases in the articles",

```

```
xlab="cancers",
ylab="coronaviruses",
keep.dendro=FALSE)
```

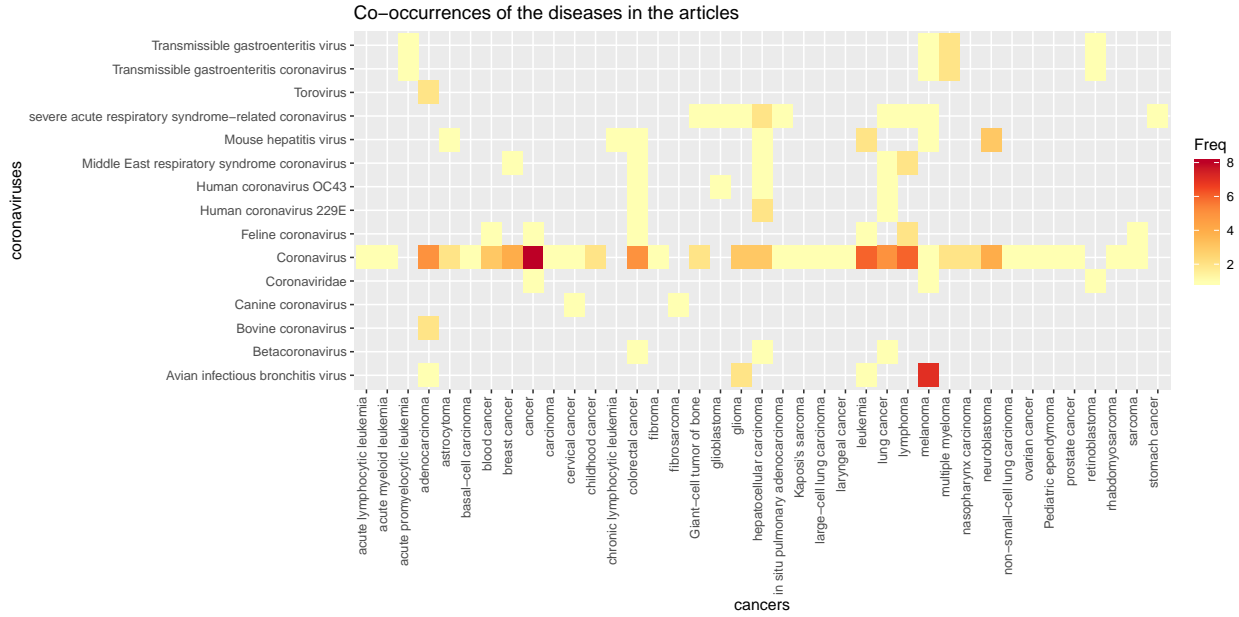


Plot heatmap with ggplot

```
library(ggplot2)
```

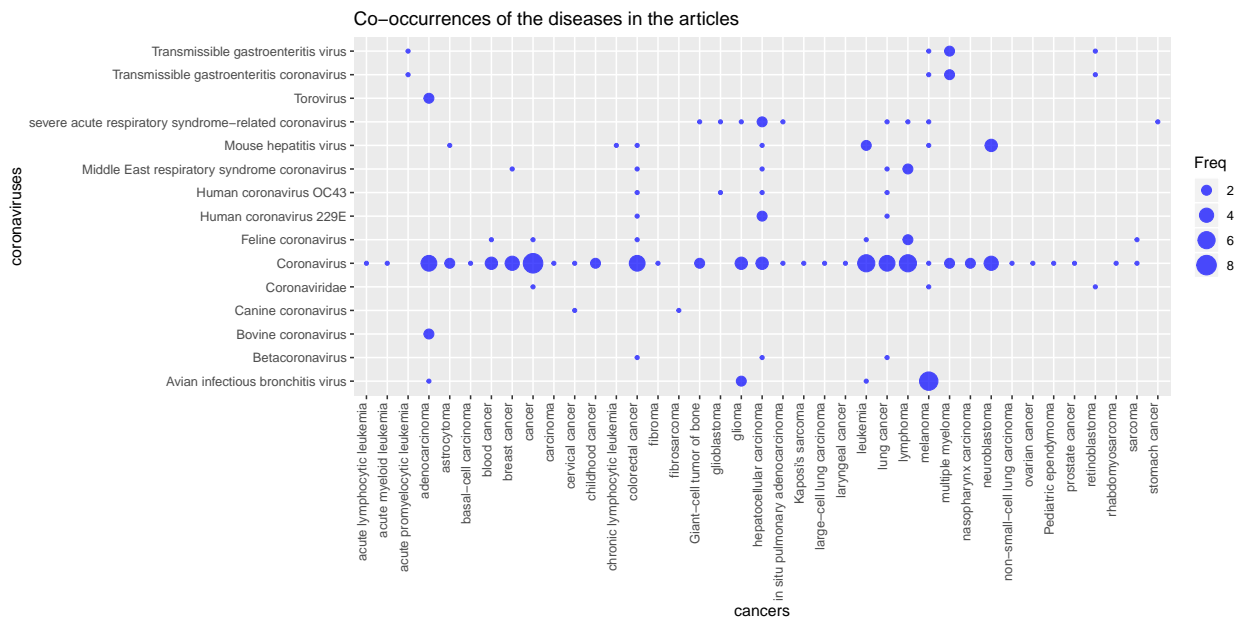
```
counts <- table(res[, c('dis2Label', 'dis1Label')])
counts <- as.data.frame(counts)
counts <- counts[counts$Freq > 0, ]

ggplot(counts, aes(x=dis1Label, y=dis2Label, fill=Freq)) +
  geom_tile() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust=0)) +
  scale_fill_distiller(palette = "YlOrRd", direction = 1) +
  ggtitle("Co-occurrences of the diseases in the articles") +
  xlab("cancers") +
  ylab("coronaviruses")
```



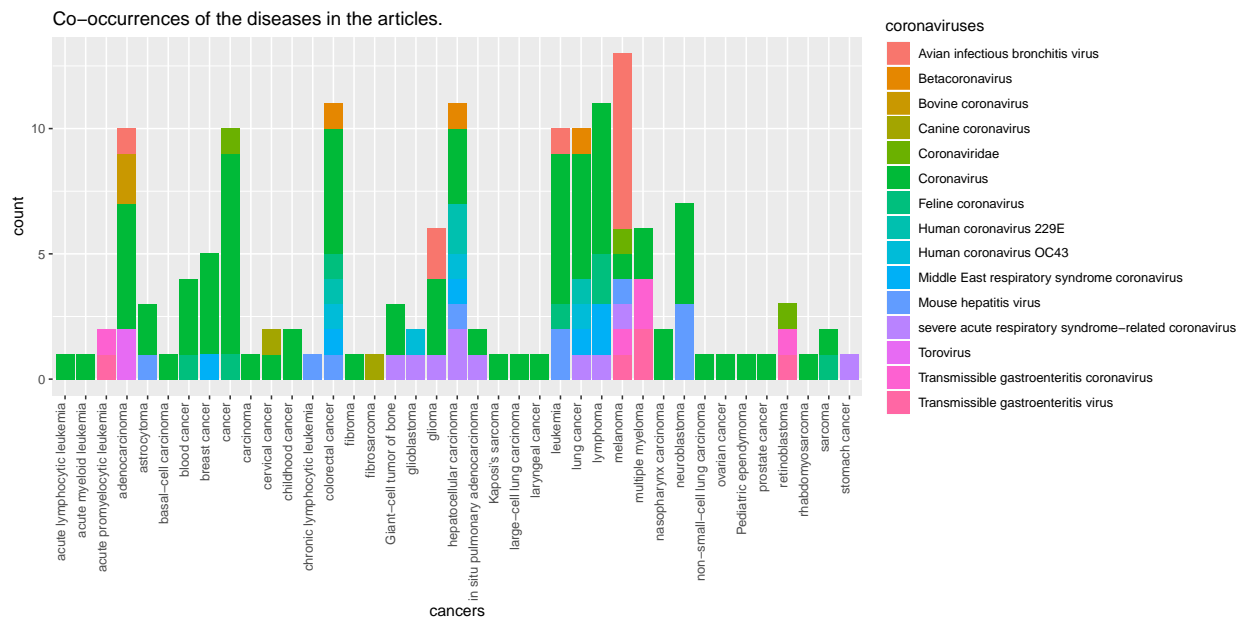
Plot bubble chart

```
# Most basic bubble plot
ggplot(counts, aes(x=dis1Label, y=dis2Label, size = Freq)) +
  geom_point(alpha=0.7, color='blue') +
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust=0)) +
  ggtitle("Co-occurrences of the diseases in the articles") +
  xlab("cancers") +
  ylab("coronaviruses")
```



Plot stacked bar chart with default colors

```
# stacked bar chart
ggplot(res,
      aes(x = dis1Label,
          fill = dis2Label)) +
geom_bar(position = "stack") +
theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust=0)) +
ggtitle("Co-occurrences of the diseases in the articles.") +
xlab("cancers") +
ylab("count") +
labs(fill = "coronaviruses")
```



Plot barchart with manual coloring

```
library(RColorBrewer)
#display.brewer.all(type="seq")
```

```
# Manually group coronaviruses in 3 groups (general(1), human(2), animal(3)) for reordering
res$group <- rep(3, nrow(res))
res$group[grep('Human', res$dis2Label)] <- 2
res$group[grep('Corona', res$dis2Label)] <- 1
res$group[grep('severe acute', res$dis2Label)] <- 2
res$group[grep('Middle', res$dis2Label)] <- 2
res$group[grep('Beta', res$dis2Label)] <- 1

# Manually choose colors for coronaviruses (shades of purple for general (1), blues for humans(2), and greens for animals(3))
greens <- brewer.pal(9, "Greens")[2:9]
purples <- brewer.pal(9, "Purples")[5:9]
oranges <- brewer.pal(9, "Oranges")[5:9]

my_colors <- c(greens[1], # avian...
```

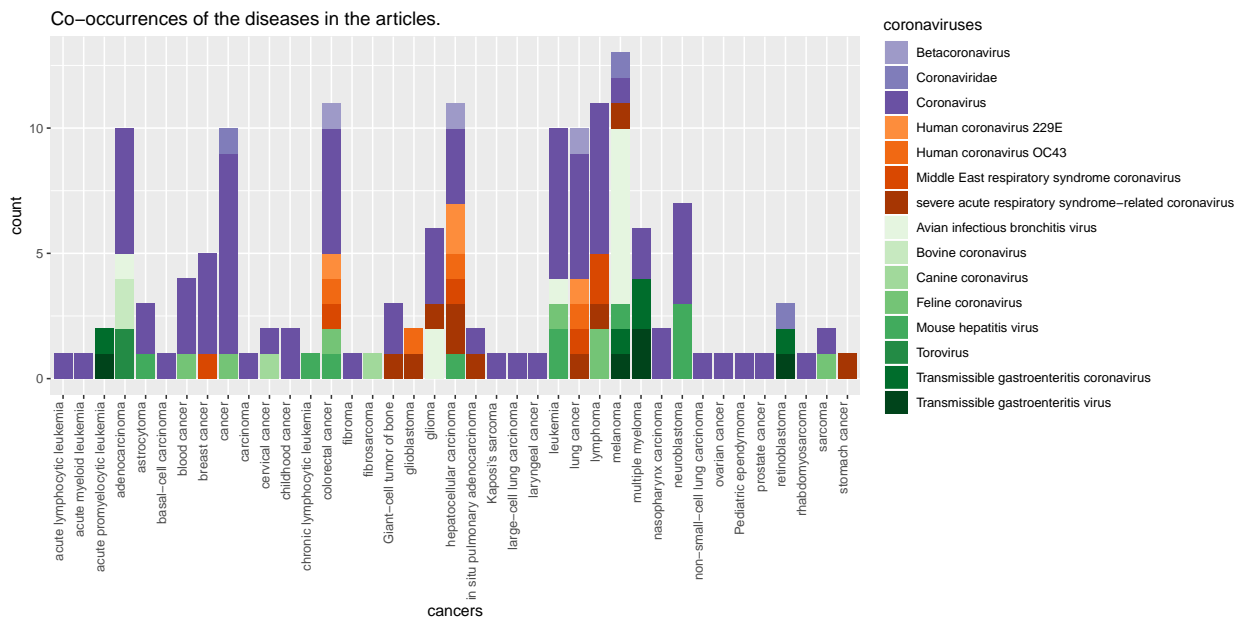
```

purples[1], # beta...
greens[2:3], # bovine..., canine...
purples[2:3], # corona..., corona...
greens[4], # feline...
oranges[1:3], # human..., human..., MERS
greens[5], # mouse...
oranges[4], # severe acute...
greens[6:9]) # torro..., gastro..., gastro...

names(my_colors) <- levels(res$dis2Label)

ggplot(res,
  aes(x = dis1Label,
      fill = reorder(dis2Label , group))) +
  geom_bar(position = "stack") +
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust=0)) +
  scale_fill_manual(name = "coronaviruses", values = my_colors) +
  ggtitle("Co-occurrences of the diseases in the articles.") +
  xlab("cancers") +
  ylab("count") +
  labs(fill = "coronaviruses")

```



Plot barplot for grouped diseases

```

group_colors = c(purples[1],
  oranges[1],
  greens[4])

ggplot(res,
  aes(x = dis1Label,
      fill = factor(group, levels = c(1, 2, 3), labels = c("General", "Human", "Animals")))) +
  geom_bar(position = "stack") +

```



```

theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust=0)) +
scale_fill_manual(name = "group", values = group_colors) +
ggtitle("Co-occurrences of the groups of coronavirus diseases in the articles.") +
xlab("cancers") +
ylab("count") +
labs(fill = "coronaviruses")

```

